

Code Documentation

1. Libraries Used

1.1 Data Manipulation

- `pandas` :
 - `pd.read_csv()` : Reads the dataset from a CSV file.
 - `pd.Series()` : Handles class distributions and counts.

1.2 Data Preprocessing

- `sklearn.preprocessing` :
 - `MinMaxScaler` : Scales feature values to a specific range (default `[0, 1]`).
 - `LabelEncoder` : Converts categorical labels into integers.

1.3 Machine Learning and Metrics

- `imblearn` :
 - `SMOTE` : Handles imbalanced datasets by oversampling the minority class.
 - `RandomUnderSampler` : An optional method to under-sample the majority class.

1.4 Neural Networks

- `tensorflow.keras` :
 - Layers: `Dense`, `Dropout`, `Input`, `LSTM`.
 - Utilities: `to_categorical` (for one-hot encoding).
 - Optimizers: Adam optimizer.
 - Initializers: `GlorotUniform`, `HeNormal`, etc.

1.5 Visualization

- `matplotlib.pyplot` : For plotting accuracy, loss curves, and SHAP visualizations.

- `seaborn`: For heatmap visualizations.

1.6 Model Explanations

- `shap`: Explains feature importance using SHAP values.

2. Data Loading and Inspection

```
dataset = pd.read_csv('/kaggle/input/combined-network-faults-  
data/combined_network_faults_data.csv')  
print(dataset.head())  
print(dataset.isnull().sum())
```

- **Purpose:**
 - Loads the dataset.
 - Displays the first few rows.
 - Checks for missing values in each column.

3. Handling Missing Data

```
dataset = dataset.dropna()  
print(dataset.isnull().sum())
```

- **Purpose:** Removes rows containing missing values and rechecks for nulls.

4. Checking Class Imbalance

```
print(dataset['0'].value_counts())
```

- **Purpose:** Displays the distribution of classes in the target column (`'0'`).

5. Handling Class Imbalance Using SMOTE

```
features = dataset.iloc[:, :-1] # All columns except the last (target).
target = dataset['0']           # Target column.

smote = SMOTE(sampling_strategy='auto')
X_resampled, y_resampled = smote.fit_resample(features, target)

print(pd.Series(y_resampled).value_counts())
```

- **Purpose:**

- Uses Synthetic Minority Oversampling Technique (SMOTE) to balance class distribution.

- **Parameters:**

- `sampling_strategy`: Specifies oversampling ratio or strategy.

6. Feature Scaling

```
scaler = MinMaxScaler()
scaled_features = scaler.fit_transform(X_resampled)
```

- **Purpose:** Scales features to a range of `[0, 1]`.

7. Label Encoding and One-Hot Encoding

```
label_encoder = LabelEncoder()
encoded_labels = label_encoder.fit_transform(y_resampled)
one_hot_labels = to_categorical(encoded_labels, num_classes=5)
```

- **Purpose:**

- Encodes target labels into integers.

- Converts integer labels into one-hot encoded vectors.

8. Reshaping Data for LSTM

```
X_resampled_resaped = scaled_features.reshape((scaled_features.shape[0], 1, scaled_features.shape[1]))
```

- **Purpose:** Reshapes the input data to fit the LSTM input format (`samples, time_steps, features`).

9. Splitting Data

```
X_train, X_test, y_train, y_test = train_test_split(X_resampled_resaped, one_hot_labels, test_size=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.2, random_state=42)
```

- **Purpose:**
 - Splits the data into training (64%), validation (16%), and testing (20%) subsets.

10. Model Building and Training

```
model = Sequential([
    LSTM(128, activation='relu', input_shape=(X_train.shape[1], X_train.shape[2])),
    Dropout(0.4),
    Dense(64, activation='relu'),
    Dense(5, activation='softmax')
])
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(X_train, y_train, validation_data=(X_val,
y_val), epochs=10, batch_size=32)
```

- **Purpose:**

- Builds an LSTM model.

- **Layers:**

- **LSTM**: Learns temporal dependencies.
- **Dropout**: Reduces overfitting by dropping nodes.
- **Dense**: Fully connected layers for classification.

- **Parameters:**

- **epochs**: Number of iterations over the entire dataset.
- **batch_size**: Number of samples per gradient update.

11. Hyperparameter Tuning

```
param_grid = {'lstm_units': [64, 128], 'dropout_rate': [0.3,
0.4], 'learning_rate': [0.001, 0.01]}
```

- **Purpose:**

- Tests different combinations of LSTM units, dropout rates, and learning rates to find the optimal configuration.

12. Visualizing Performance

```
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation A
ccuracy')
```

- **Purpose:** Plots accuracy and loss curves for each training configuration.

13. Feature Importance (SHAP)

```
explainer = shap.KernelExplainer(predict_fn, X_val_flat[:10
0])
shap_values = explainer.shap_values(X_val_flat[:10])
shap.summary_plot(shap_values, X_val_flat[:10], feature_names
=feature_names)
```

- **Purpose:** Explains feature contributions using SHAP values.

14. Weight Initialization Techniques

```
initializers = {
    'glorot_uniform': GlorotUniform(),
    'he_normal': HeNormal(),
    'orthogonal': Orthogonal(),
    'random_normal': RandomNormal(mean=0.0, stddev=0.05),
    'lecun_uniform': LecunUniform()
}
```

- **Purpose:** Tests different weight initializations for better model convergence.

Explanation of Weight Initialization Techniques

Weight initialization is crucial to deep learning, as it determines how well a model converges during training. These techniques set the initial values of a model's weights before training begins, ensuring gradients flow properly through the network and preventing issues like vanishing or exploding gradients.

Here's a detailed explanation of each initializer:

1. **glorot_uniform** (Xavier Initialization)

- **Definition:** The weights are initialized from a uniform distribution with bounds where:

$$\left[-\sqrt{\frac{6}{n_{\text{in}}+n_{\text{out}}}}, \sqrt{\frac{6}{n_{\text{in}}+n_{\text{out}}}}\right],$$

- `ninn_{\text{in}}`: Number of input units.
 - `noutn_{\text{out}}`: Number of output units.
 - **Purpose:** Balances the variance of weights across layers, helping the gradient propagate effectively during backpropagation.
 - **Use Case:** Suitable for activations like `tanh` or `sigmoid`.
 - **Advantages:** Works well for shallow and deep networks, reducing the risk of vanishing/exploding gradients.
-

2. `he_normal` (He Initialization)

- **Definition:** The weights are initialized from a normal distribution with a mean of 0 and a standard deviation of $2n_{\text{in}} \sqrt{\frac{2}{n_{\text{in}}}}$, where `ninn_{\text{in}}` is the number of input units.
 - **Purpose:** Addresses the limitations of Xavier initialization, especially for rectified linear unit (ReLU) activations, by considering their asymmetric behavior (outputs non-zero only for positive inputs).
 - **Use Case:** Ideal for layers with ReLU or its variants (e.g., LeakyReLU).
 - **Advantages:** Ensures that the variance of weights remains appropriate for ReLU activations.
-

3. `orthogonal`

- **Definition:** The weights are initialized to form an orthogonal matrix, where the dot product of any two rows/columns is zero.
 - **Purpose:** Maintains the independence of weights, ensuring that information flows through the network without interference.
 - **Use Case:** Commonly used for RNNs, LSTMs, and GRUs, where preserving orthogonality is important to avoid exploding or vanishing gradients over long sequences.
 - **Advantages:** Helps stabilize gradients in recurrent architectures.
-

4. `random_normal`

- **Definition:** The weights are initialized from a normal distribution with a specified mean (default 0.0) and standard deviation (default 0.05).
- **Purpose:** Provides control over the randomness of weight initialization.
- **Use Case:** Useful for custom experiments where specific ranges or distributions are required.
- **Advantages:**
 - Flexibility in defining the mean and variance.
 - Simpler but less sophisticated compared to Xavier or He initialization.

5. `lecun_uniform`

- **Definition:** The weights are initialized from a uniform distribution with bounds $[-\sqrt{\frac{3}{n_{\text{in}}}}, \sqrt{\frac{3}{n_{\text{in}}}}]$, where n_{in} is the number of input units.
- **Purpose:** Specifically designed for use with `selu` (Scaled Exponential Linear Unit) activations.
- **Use Case:** Used in networks where self-normalizing properties are desired, particularly for deep architectures.
- **Advantages:** Works well with `selu` activations to achieve self-normalizing behavior in deep networks.

Summary Table of Initializers

Initializer	Distribution	Formula/Properties	Best For
<code>glorot_uniform</code>	Uniform	$[-\sqrt{\frac{6}{n_{\text{in}}+n_{\text{out}}}}, \sqrt{\frac{6}{n_{\text{in}}+n_{\text{out}}}}]$	<code>tanh</code> , <code>sigmoid</code> activations
<code>he_normal</code>	Normal	Mean = 0, StdDev = $\sqrt{\frac{2}{n_{\text{in}}}}$	<code>ReLU</code> and variants
<code>orthogonal</code>	Matrix	Orthogonal weight matrix	RNNs, LSTMs, GRUs
<code>random_normal</code>	Normal	Mean and StdDev are configurable	Custom experiments
<code>lecun_uniform</code>	Uniform	$[-\sqrt{\frac{3}{n_{\text{in}}}}, \sqrt{\frac{3}{n_{\text{in}}}}]$	<code>selu</code> activations

Choosing the right initializer depends on the network architecture and activation functions used. For most general cases:

- Use `he_normal` with ReLU-like activations.
- Use `glorot_uniform` with sigmoid/tanh activations.
- Use `orthogonal` for recurrent networks.
- Experiment with `random_normal` and `lecun_uniform` in specific scenarios.

15. Prediction on Unseen Data

```
unseen_data_scaled = scaler.transform(unseen_data)
unseen_data_resaped = unseen_data_scaled.reshape((unseen_data_scaled.shape[0], 1, unseen_data_scaled.shape[1]))
prediction = model.predict(unseen_data_resaped)
predicted_class = np.argmax(prediction)
```

- **Purpose:** Predicts the class of new input data.

16. Key Parameters

Data Preprocessing

Parameter	Description
<code>sampling_strategy</code>	Ratio for SMOTE oversampling.

Model Training

Parameter	Description
<code>lstm_units</code>	Number of units in the LSTM layer.
<code>dropout_rate</code>	Fraction of units to drop in dropout layer.
<code>learning_rate</code>	Optimizer step size.
<code>epochs</code>	Number of complete passes over the dataset.
<code>batch_size</code>	Number of samples per gradient update.

